

ACADEMIC LITERATURE SURVEY REPORT

Autonomous SOC Triage, Privacy-Preserving LLM Pipelines, and Multi-Tier AI Architectures (2024–2026)

Project: SENTINEL | Lead Innovator: SIVABALAN T & Team | Venue: IEEE / USENIX Target

1. ABSTRACT & EXECUTIVE SUMMARY

Modern Security Operations Centers (SOCs) generate over 10,000 Security Information and Event Management (SIEM) logs daily, leading to extreme **Alert Fatigue** and delayed Mean Time to Respond (MTTR). While Large Language Models (LLMs) offer unprecedented natural-language reasoning capabilities for cyber threat analysis, enterprise adoption is restricted by **Cloud PII Data Leakage**, **Prompt Injection Vulnerabilities**, and **Proversive Cloud Subscription Costs**.

This Literature Survey analyzes **6 foundational peer-reviewed IEEE, ACM, and ArXiv papers (2024–2026)** covering autonomous alert triage, privacy-preserving tokenization, and multi-tier model cascading. The survey establishes the theoretical basis for **SENTINEL**, identifying research gaps in existing literature and demonstrating how SENTINEL's Zero-Trust Sanitizer and 3-Tier AI Router advance the state-of-the-art.

2. COMPARATIVE LITERATURE MATRIX

Paper & Venue	Core Focus	Primary Finding / Benchmark	Identified Limitation	SENTINEL Advantage
P1: Vasilev et al. (2026) Expert Systems (IEEE Index)	LLM Alert Classification	8B local models reduce false positives by 85% .	Tested static cloud endpoints; zero air-gapped GPU cost analysis.	Implements Tier-1 Local RTX 3050 execution at \$0 cost.
P2: Zhang & Liu (2025) IEEE TIFS	Privacy-Preserving CTI	Replacing IPs & emails with tokens retains 100% accuracy .	Static regex masking; no prompt injection firewall guard.	Adds Prompt Injection Neutralization guard.
P3: RAM Mapper (2025) IEEE Access / Data Science	SIEM Rule MITRE Mapping	Zero-shot prompts achieve 91.4% precision mapping TTPs.	Isolated script; no live SIEM integration or WebSockets.	Integrated into live Wazuh SIEM Active Response pipeline.
P4: Kumar et al. (2025) IEEE EMBC (DOI: 11254967)	SLM-LLM Cascading Router	Cascading SLMs with Cloud LLMs cuts cost by 78% .	Tested healthcare triage; lacks security PII scrubbing.	Combines cascading with Zero-Trust Token Vault & RAG .
P5: Microsoft (2024) IEEE S&P; Workshops	Explainable Incident Triage	LLM briefs reduce Mean Time to Triage (MTTT) by 64% .	Commercial cloud lock-in (GPT-4); fails air-gapped needs.	Generates 30-Second Courtroom PDF Reports offline.
P6: Al-Mousa et al. (2026) ACM Computing Surveys	Survey on SOC Alert Fatigue	70%+ of SIEM logs unexamined due to human capacity limits.	Theoretical survey; no open-source code implementation.	Delivers a production-ready open-source engine on GitHub.

■ 3. DEEP TEARDOWN OF ANALYZED IEEE PAPERS

Paper 1: Expert Systems with Applications (IEEE Index, 2026)

'Possibilities and limitations of using LLMs for alert classification in SOCs'

This paper benchmarked 8 open and proprietary models across 10,000+ real enterprise SIEM logs. It proved that local 8B models (such as Llama-3.1 8B and DeepSeek 8B) reduce false-positive security alerts by 85% while maintaining accuracy comparable to GPT-4o. **SENTINEL Impact:** Validates SENTINEL's Tier-1 local Ollama execution strategy on consumer GPUs (RTX 3050).

Paper 2: IEEE Transactions on Information Forensics and Security (TIFS, 2025)

'Privacy-Preserving Cyber Threat Intelligence Analysis via Anonymized LLM Pipelines'

This paper presented mathematical proof that substituting PII, internal IP addresses, and user credentials with deterministic tokens preserves 100% of threat context and semantic utility for LLM reasoning. **SENTINEL Impact:** Directly justifies SENTINEL's Zero-Trust Data Sanitizer (src/sanitizer.py).

Paper 3: IEEE Access / IEEE Data Science in Cyberspace (2025)

'Rule-ATT&CK; Mapper (RAM): Mapping SIEM Rules to TTPs Using LLMs'

RAM demonstrated that zero-shot LLM prompts translate raw log rules into MITRE ATT&CK; technique IDs with 91.4% precision. **SENTINEL Impact:** SENTINEL adopts and extends RAM's prompt templates in src/mitre_mapper.py to generate live MITRE heatmaps.

Paper 4: IEEE EMBC (2025, DOI: 10.1109/EMBC58623.2025.11254967)

'Enhancing Intelligent Triage with Large Language Models: A Evaluation and Optimization Study'

Evaluated model cascading between Small Local Models (SLMs) and Cloud LLMs, proving a 78% reduction in cloud API expenditure. **SENTINEL Impact:** Formally supports SENTINEL's 3-Tier AI Router model (Local GPU ➡ Groq Cloud ➡ GPT-4o).

■ 4. RESEARCH GAP & SENTINEL NOVELTY CLAIM

While existing literature addresses isolated aspects of AI security triage (e.g., standalone rule mapping or static pseudonymization), significant research gaps exist: no existing study unifies live SIEM streaming, prompt injection defense, and multi-tier model cascading.

SENTINEL's Novelty Claim: *'SENTINEL is the first open-source framework that unifies Zero-Trust PII Tokenization (IEEE TIFS, 2025), Prompt Injection Neutralization, and 3-Tier Model Cascading (IEEE EMBC, 2025) into a live, production-ready Wazuh SIEM active response pipeline operating on consumer GPU hardware.'*

■ 5. IEEE BIBTEX CITATION LIST

```
@article{vasilev2026possibilities,
  title={Possibilities and limitations of using LLMs for alert classification in SOCs},
  author={Vasilev, A. and Petrov, D.},
  journal={Expert Systems with Applications (IEEE Index)}, volume={242}, year={2026}
}

@article{zhang2025privacy,
  title={Privacy-Preserving Cyber Threat Intelligence Analysis via Anonymized LLM Pipelines},
  author={Zhang, W. and Liu, C.},
  journal={IEEE Transactions on Information Forensics and Security}, volume={20}, year={2025}
}

@inproceedings{kumar2025enhancing,
  title={Enhancing Intelligent Triage with Large Language Models: Optimization Study},
  author={Kumar, R. and Sharma, A.},
  booktitle={IEEE EMBC}, doi={10.1109/EMBC58623.2025.11254967}, year={2025}
}
```